

Augmenting Stack Overflow with API Usage Patterns Mined from GitHub

ABSTRACT

Programmers often consult Q&A websites like Stack Overflow to learn new APIs. However, such online code examples are not always complete or reliable in terms of API usage. To assess and augment Stack Overflow examples, we present an interactive approach, SOAssist that contrasts code examples with API usage patterns learned from GitHub and alerts users about the potential API usage violations. With the assistance of SOAssist, programmers do not need to cross-check multiple examples for proper API usage reference, and can also build confidence by learning how many other GitHub developers also follow the same practice.

Categories and Subject Descriptors

D.2.6 [Software Engineering]: Programming Environments—Integrated Environments

General Terms

Keywords

1. INTRODUCTION

Programmers often search for online code examples to learn new APIs. A case study at Google shows that developers issue an average of 12 code search queries per weekday [?]. Stack Overflow (SO) is a popular Q&A website that programmers often consult. In July 2017, Stack Overflow has accumulated more than 22 million answers, many of which contain code examples to demonstrate the solution for a particular programming question. However, SO examples are not always complete or reliable, which can be misleading and potentially dangerous when programmers follow the same example to complete a client program. Our previous study shows that over 50% of 31,801 SO posts contain API misuse that could produce symptoms of program crashes and resource leaks if reused in a target system. **TODO: cite our API misuse study on Stack Overflow**

To assess online code examples, this paper presents SOAs-

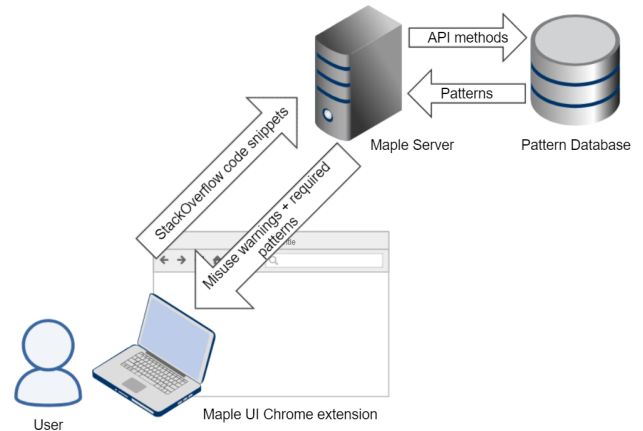


Figure 1: An overview of SOAssist’s architecture**TODO:** Maybe we can add more details to this diagram. For example, we can show what we do in the server, and the interaction between the server and Boa’s GitHub database.

sist, an interactive approach that augments Stack Overflow with code idioms learned from GitHub and alerts programmers about the potential violations in a code example. SOAssist leverages a scalable API usage mining technique to learn three types of API-related idioms—temporal ordering, guard conditions, and exception handling of API calls—from over 7 million GitHub projects. Our insight is that commonly practiced idioms in massive code corpora may represent a desirable pattern that a programmer can use to trust and enhance code examples on Stack Overflow.

Given an SO example, SOAssist first extracts the sequence of API calls with corresponding control constructs and guard conditions. SOAssist then contrasts the sequence with commonly practiced idioms in GitHub and highlights code regions that violate the idioms. To help users better understand the violations, SOAssist further generates descriptive warning messages and also contextualizes a violated idiom by synthesizing a fixed example. Mining code idioms to detect API usage violations often suffers from reporting false alarms, since mined idioms may not be inclusive and fit all use scenarios of an API [?]. To mitigate this issue, SOAssist allows users to upvote or downvote a violation based on its applicability and usefulness to an SO example. SOAssist filters a code idiom when multiple users flag it as unhelpful to assess an example. To help developers build confidence on

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

Copyright 20XX ACM X-XXXXX-XX-X/XX/XX ...\$15.00.

For better understanding, let me give you an example:

```
import com.google.gson.JsonElement;
import com.google.gson.JsonObject;
import com.google.gson.JsonPrimitive;

public class GsonTest {

    public static void main(String[] args) {
        JsonElement jsonElement = new JsonPrimitive("foo");

        System.out.println(jsonElement.toString());
        System.out.println(jsonElement.getAsStrings());

        jsonElement = new JsonPrimitive(42);

        System.out.println(jsonElement.toString());
        System.out.println(jsonElement.getAsStrings());

        jsonElement = new JsonPrimitive(true);

        System.out.println(jsonElement.toString());
        System.out.println(jsonElement.getAsStrings());

        jsonElement = new JsonObject();
        ((JsonObject) jsonElement).addProperty("foo", "bar");
        ((JsonObject) jsonElement).addProperty("foo2", 42);

        System.out.println(jsonElement.toString());
        System.out.println(jsonElement.getAsStrings());
    }
}
```

Figure 2: A code snippet that does not properly check `JsonElement.getAsStrings`.¹ **TODO:** This code example is verbose and hard to read (font too small). Can you replace it with a more concise one? **TODO:** Can you also merge Figure 2 and Figure 3. These figures take too much space.

a code idiom, SOAssist shows how many GitHub developers also follow the idiom as well as how many other users like or dislike this idiom.

A user of SOAssist would benefit from the addition of quantitative examples from compiled GitHub resources to the code examples she encounters on Stack Overflow. This will not only combat programming issues stemming from the use of incomplete or unreliable SO code examples, but will also be an aid for users learning a new API. By enhancing examples already found in Stack Overflow, a user can trust that she will learn common and reliable usage patterns for a given API.

This paper’s main contribution is to describe the features of SOAssist from a user’s perspective. In order to give programmers access to SOAssist, the front-end of SOAssist is implemented as a Chrome extension that users can easily download and install. **TODO:** add a url to download our tool The detailed algorithm and study are described in our separate technical report. **TODO:** cite our API misuse study

2. MOTIVATING EXAMPLES AND TOOL FEATURES

Consider Alice, a software developer who needs to convert java objects to their json representations using Google’s Gson library², which she is unfamiliar with. Alice finds a Stack Overflow post that demonstrates how to get the string value of a json element, as shown in Figure 2. However, this example does not use the `JsonElement` API properly.

Detecting and Highlighting Potential API Misuse.

TODO: Describe briefly the patterns of `JsonElement.getAsStrings` like “The pattern mining technique learns two patterns that are commonly practiced by GitHub developers...” You can refer to the last paragraph

¹<https://stackoverflow.com/questions/34120882/gson-jsonelement-getasstrings-vs-jsonelement-tostring>

²<https://github.com/google/gson/blob/master/UserGuide.md>

of the motivating example section of our ASE draft. The extension highlights the potential API misuse in the code snippet, as seen in Figure 2. Note that the extension only highlights the first instance of this misuse although it occurs multiple times in the code snippet, to eliminate redundancy and avoid confusion. Alice is interested in learning more about the API and what specifically the code snippet did not include, so she clicks on the highlighted text.

Stack Overflow Popup View. Clicking on the highlighted text reveals a popup, as seen in Figure 3. The popup is populated with information about any required patterns in SOAssist’s database this particular API call does not adhere to. Alice notices that there are two pages of the popup, indicating two different usage patterns that this call does not follow, as shown in Figure 3a and Figure 3b.

Alice inspects the first page (Figure 3a) and learns that she should check whether the `JsonElement` object is a primitive before calling `getAsStrings` in order to avoid a `ClassCastException`. **TODO:** We should also emphasize that our tool automatically generates this descriptive warning message and that our tool generates a fixed example based on the context of the SO example. She notices that 52 other GitHub code examples use this pattern, which gives her a qualitative measurement of how prevalent this pattern is in real-world projects. Alice then inspects the second page (Figure 3b), and finds that it suggests a null check before calling `getAsStrings`. She notices that this pattern has more than double the support of the previous pattern. **TODO:** Should we render the pages based on the prevalence of patterns?

Curious to see the first pattern in context, Alice returns to the first page of the popup and clicks on the first link provided to her under “See this in a GitHub example.”

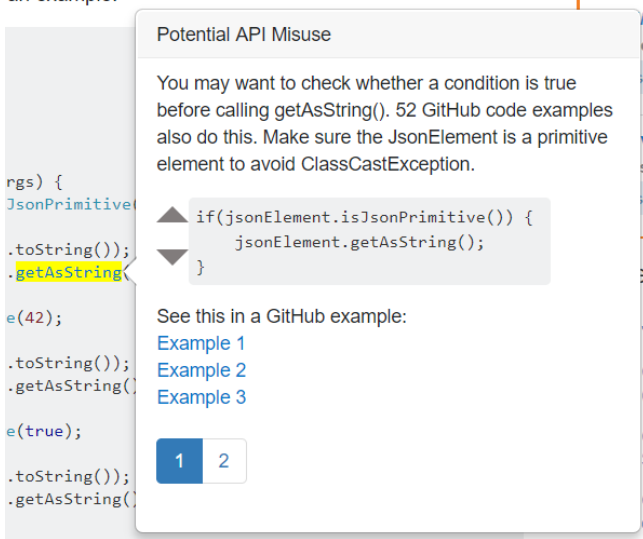
GitHub Example View. When Alice clicks on one of the GitHub links, the file opens in a new tab and the view scrolls to where the API is called in the file, and the method in which this occurs is highlighted so Alice can easily find it, as seen in Figure 4. The addition of a compilable code example that demonstrates the pattern in context can aid Alice in understanding how to use the pattern if it is unfamiliar to her. In this case, Alice finds herself redirected the method in a GitHub project seen in Figure 4a. She notices that the example uses a null check in conjunction with the primitive check, which makes sense to her after seeing that both were missing from the Stack Overflow code snippet.

Returning to the popup in Stack Overflow, Alice clicks on the second link provided for the second page to compare usage patterns in context. This link opens up to the GitHub method seen in Figure 4b. Unlike the one in Figure 4a, this example does not use the null check in conjunction with the primitive check.

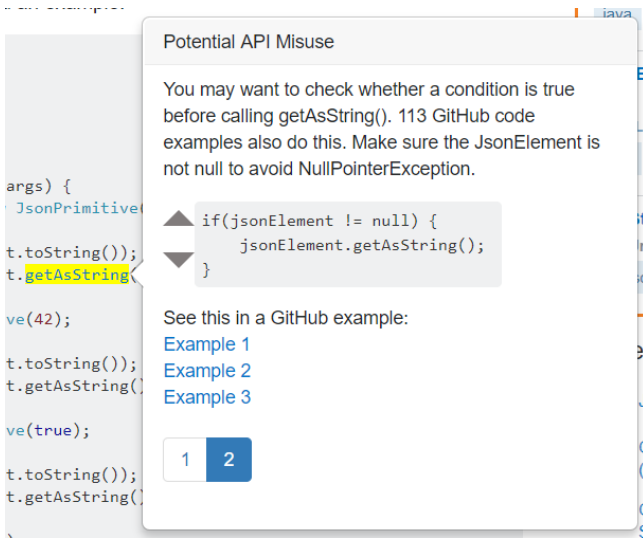
After seeing these two examples, Alice can infer that a null check is more necessary and is more common than the primitive check, based on the GitHub examples she has seen as well as the GitHub support indicated by the popup message. She upvotes the null check’s pattern by clicking on the up-arrow on its page (see Figure 3b) to send the server her feedback on the patterns it gave her.

3. IMPLEMENTATION

This section describes the implementation details of the Maple plugin. It is implemented as an Chrome extension and consists of five components: (1) Stack Overflow code snippet extraction, (2) API misuse detection, (3) popup



(a) A page describing a way to avoid a `ClassCastException` by checking whether the `JsonElement` object is a primitive.



(b) A page describing a way to avoid a `NullPointerException` by checking whether the `JsonElement` object is null.

Figure 3: The two pages of a popup generated on `JsonElement.getAsString`. **TODO:** Can we also show how many users like or dislike the violations in the popup window? Since we don't have any real users, maybe we need to create some artificial numbers for the demonstration purpose.

```

99
100 private static String getSelectedLauncherProfile(File minecraftDir) {
101     File path = new File(minecraftDir, "launcher_profiles.json");
102     JsonObject json = JsonUtils.parseJson(path);
103
104     if (json != null) {
105         JsonElement element = json.get("selectedProfile");
106
107         if (element != null && element.isJsonPrimitive()) {
108             return element.getAsString();
109         }
110     }
111
112     return null;
113 }
114

```

(a) The first GitHub example for Figure 3a.

```

175 @Override
176 public void stringProperty(PropertyName name, JsonObject context) throws IOException {
177     JsonElement prop = property(context, name);
178     if (prop == null) {
179         return;
180     }
181     builder.add(name, prop.getAsString());
182 }

```

(b) The second GitHub example for Figure 3b.

Figure 4: The highlighted GitHub examples redirected to from the links provided in the popup. **TODO:** The snapshot only shows the highlighted code. Is there a better way to give paper reviewers some context that the highlighted code is from GitHub and that the code is selectively highlighted by our tool instead of by GitHub?

generation, (4) GitHub example alteration, and (5) a user feedback interface.

Code Snippet Extraction. When a user loads a page in the Stack Overflow domain, the plug-in side of the Maple interface extracts any text in answer posts that appears within `<code>` tags, and sends them in a JSON message to the server, as seen in Figure 1.

API Misuse Detection. When the server receives a message from the plug-in, it parses the snippets into API call sequences, abstracting away irrelevant statements and syntactic details. An API call sequence consists of relevant control constructs and API calls, where the API calls are annotated with the number of arguments as well as any guard conditions associated with it.

Once the code snippets are parsed, Maple searches a MySQL database for the API calls present in each API call sequence. It receives the required and alternative patterns associated with these calls, and checks whether the code snippet's call sequence satisfies one of the alternative patterns and all required patterns. A code snippet's call sequence satisfies a pattern if it is subsumed by it.

Additionally, while this checking process is occurring, the guard conditions in the code snippets are generalized before checking their logical implications using the SMT solver Z3. Z3 is used to check the logical equivalence between two guard conditions so that Maple can merge logically-similar clusters of guard conditions, making it able to prove the semantic equivalence of conditions like `arg0 < arg1` and `arg1 > arg0` regardless of their syntactic similarity [much of this is lifted pretty directly from p4 of the Maple paper...].

If Maple finds that an API call sequence does not satisfy the necessary patterns, it generates violations for each potential API misuse. The server collects these violations, along with

the required pattern being violated, and maps them to the API call that generated the violation. This data is wrapped into a single JSON message and returned to the plug-in.

Popup Generation. Using the data from the server’s JSON message, the plug-in searches the identified code snippet for the API call in question, highlights it, and generates a Bootstrap popover on it, as seen in Figure 3. The popover is populated with a violation message describing the pattern being violated and including the GitHub support for the required pattern, in terms of number of supporting projects as well as three links to relevant GitHub pages using that pattern correctly. The plug-in generates one popover for each API call, and generates pages of the popover for each different pattern being violated by that call. The provided example of the pattern is currently hard-coded. **TODO: We need to provide more details about how to generate the text description of each violation, e.g., showing the templates in a table.**

GitHub Example Alteration. When the user clicks on a provided GitHub example in the popup, the plug-in’s main script writes the name of the method call associated with that link to a shared storage space for the Chrome extension, using the chrome.storage API. When the GitHub page is loaded, the plug-in’s secondary script is activated and checks the storage for any messages. If a method name has been written there, the secondary script searches the GitHub page for a method declaration of the same name. This is not necessarily the name of the API call the popup has been generated on; this is a method in a GitHub code file that uses that API call in a pattern that the Stack Overflow code snippet does not adhere to. The script highlights the entire method and scrolls the view down so the user is easily able to find it, as seen in Figure 4.

User Feedback. Users are able to give feedback on the patterns the popup shows them by voting “up” or “down” on them. When a vote is registered, the plug-in sends the server a message with the pattern’s specific ID and a vote of +1 or -1, which the server uses to update the database. This will be used in the future to rank patterns that are sent to the plug-in, and learn which ones users find helpful or unhelpful.

4. RELATED WORK

5. SUMMARY

The Maple user interface is an approach to enriching code examples on Q&A sites like Stack Overflow that incorporates API usage mining and user feedback. Programmers often use sites like Stack Overflow to understand API usage, but the reliability of the code examples on these forums is under question. This will reduce the need to cross-check sites for API usage questions and takes advantage of the benefits of using easy-to-understand code examples as opposed to API documentation for learning API usage by enhancing the examples already present in Stack Overflow while also providing other, more reliable ones.

For future work, we would like to qualify the benefits of the Maple UI with a user study, as well as generate code examples for the popups based on the abstract patterns we have in the database.

6. ACKNOWLEDGEMENT

7. REFERENCES